

Davirson Novoa Ramírez

Finance Data Analyst · Financial BI Analyst · Analytics Engineer

I find the variance, I trace it to root cause, I leave the control running itself, and I explain it where the decision gets made. All four, with no translator between the business and the data.

Bogotá, Colombia · GMT-5 · Fully remote · English B2 · Portuguese A2

davirson@davirson.com · linkedin.com/in/davirson-novoa · github.com/DavinsonR · davirson.com

Profile

Economist with three years inside corporate finance — treasury, billing and FP&A — building the data the business runs on rather than waiting for someone to hand it over: Python, SQL and Power BI applied to operations across 15+ countries in the Americas and the Eastern Hemisphere. Alongside that I run a production PostgreSQL warehouse in medallion architecture, with dbt transformations, 89 automated quality tests and a Power BI semantic model, refreshed daily with no manual step, plus a credit decisioning system over 95 million records with ten model-governance gates. All of it verifiable: every figure on my site links to the repository that produces it. I am looking for a remote role where financial judgment and data engineering are paid as one capability, not two halves.

The crossover role

This is not a career change; it is pricing what I already do. In every finance role I ended up building the same thing — code and data — because the report the business needed did not exist and nobody else was going to build it. An FX mismatch analysis in Python as an intern. Automation that gave every treasury analyst about 10 hours a month back — close to 60 a month across the team. Power BI models that turned a monthly close into a decision. The edge is not knowing Python; plenty of people know Python. It is knowing which question is worth answering before writing the first line. A Finance Data Analyst is neither a data analyst who picked up finance nor a finance person who picked up code — it is the one who needs no translator between them. Hiring that person is hiring two profiles in one, without the handoff meeting between them.

Experience

Neoris EPAM

Business Consultant, FP&A

Bogotá, Colombia · remote

Mar 2026 — present

- Designed and led financial management systems for SG&A close and forecast across 15+ countries in the Americas and the Eastern Hemisphere (India, Hungary, Spain, Czech Republic).
- Managed monthly close and company-wide forecast confirmation, analyzing variances against plan and prior forecast.
- Tracked FX variance in constant currency and per-currency inflation.
- Built FP&A reporting with advanced Excel, Power Query, Power BI, MicroStrategy, JD Edwards and SAP.

SLB

Internal progression: intern to specialist in 26 months.

Bogotá, Colombia · hybrid

Billing Specialist (Argentina & Brazil)

Dec 2024 — Mar 2026

- Executed revenue recognition in SAP under Sarbanes-Oxley (SOX) requirements.
- Built Power BI models and visualizations for revenue and trend analysis.
- Performed internal audits on integrated technology projects, ensuring traceability and compliance.

Treasury Analyst

Oct 2024 — Dec 2024

- Automated bank reconciliation in SAP, improving accuracy and efficiency.
- Cut roughly 10 hours of manual reporting per analyst per month — about 60 hours a month back to the team — by optimizing and automating the process.
- Implemented automated workflows with Power Automate.

Treasury Intern

Jan 2024 — Jun 2024

- Built a comprehensive FX mismatch analysis for Latin American markets in Python, identifying financial risks and hedging opportunities.
- Created Power BI dashboards for cash-flow projections across 15+ countries.
- Automated bank transaction reporting across 10+ regional operations.

LEE Javeriana

Research Assistant

remote

Jul 2023 — Dec 2023

- Analyzed Colombian education-system data with a focus on higher education.
- Built Power BI dashboards to visualize research findings.
- Contributed to impact-evaluation studies through statistical analysis.

Solidariamente (volunteer)

Social Data Analyst

remote

Jan 2023 — Jul 2023

- Delivered analytical tools and insights for NGOs in multicultural environments.
- Collaborated with international teams on data-driven social-impact initiatives.

Production projects

Open source and checkable. None of this is a course exercise.

credit-risk-mlops

2026 — in production

Credit decisioning system with model governance — design and build

github.com/DavinsonR/credit-risk-mlops

- Production model at AUC 0.7005 over 1.96M SBA 7(a) loans and 62.4M HMDA applications: +0.0311 over an interpretable WoE scorecard, with calibration error 0.0107. The event study runs separately, over 93.4M applications.
- The first reading was 0.9461, and it was not a model but a leak: TermInMonths is overwritten when a loan is liquidated, so the field carried the outcome. Removing it drops the ablation to 0.6621 — and that is the number published.
- Ten promotion gates judged by exit code, one of them designed to block my own model when it misses the threshold.
- Monitoring caught the source changing its vocabulary without notice: a field stopped meaning the same thing between vintages, and nobody would have known without the control.
- The production ONNX artifact scores in the reader's own browser — no server, and no data leaves the page.

Stack: Python · LightGBM · scikit-learn · optbinning (WoE) · MLflow · ONNX · DuckDB · pandera · Docker · Power BI

JARVIS — personal tracking app

2026 — open demo

Own product — designed, built and operated

[public demo, no account](#)

- Multi-user model on Postgres with a row policy on 34 tables: the conversion required rebuilding 22 primary keys that identified rows without the user, and pushing the user inside the foreign keys so isolation did not rest on the policy alone.
- Eight CI gates judged by exit code: types, format, lint, 526 tests, build, dependency audit, a guard that hunts personal data inside the built artifact, and a smoke test requiring the anonymous role to be denied across all 55 tables and views.
- A public demo running the five real screens over 120 seed-generated days that never queries the database; a test walks the runtime import graph and fails if any reachable module gets to the data client.
- Day logging in one form and one button, with eight independent writes, an on-device draft and send queue to survive losing signal, and a purpose-built number reader because a browser number field was storing expenses divided by a thousand.
- Hand-written SVG charts, no library, with contrast validated in light and dark and states separated by shape as well as color.

Stack: Next.js 16 · React 19 · TypeScript · Postgres (Supabase) · RLS · Tailwind · PWA · Vitest

market-data-medallion

2026 — in operation

Market data platform — designed, built and operated

github.com/DavinsonR/market-data-medallion

- Watermark-based incremental ingestion from 4 sources (Tiingo, Coinbase, Kraken, FX) into a three-layer PostgreSQL warehouse: 48 assets and more than 58,000 candles, idempotent and resumable.
- 89 automated dbt data-quality tests plus 171 Python unit tests; the pipeline fails before publishing bad data, not after.
- No-look-ahead backtesting engine with 70/30 out-of-sample validation over 1,392 strategy variants: fewer than one in six in-sample winners survived the blind window.
- FX decomposition for Latin American ADRs, separating company performance from the currency move through the identity $(1+r_USD) \times (1+r_FX) = (1+r_local)$.
- Daily orchestration on GitHub Actions with a rate-limit circuit breaker, running on zero-cost infrastructure.

Stack: PostgreSQL · Medallion architecture · dbt · Python · Prefect · GitHub Actions · Supabase · Power BI (TMDL / PBIP) · Next.js

Research

My own research, built in public on open data.

Financial inclusion and regional growth in Colombia

2026 — thesis filed

M.Sc. in Economics thesis, Javeriana — reproducible research end to end

[repository on GitHub](#)

- Purpose-built open-data downloader with a verifiable manifest: nineteen sources (financial supervisor, statistics office, ICT and education ministries, national map), row counts verified against the source, metadata-driven typing and a sha256 per file.
- Dimensional warehouse in dbt on DuckDB: department, municipality and period dimensions, financial inclusion in long form by product block, and tests that reconcile each department total with the sum of its municipalities.
- Geographic keys resolved to the municipal code with 100% coverage across the 34 quarterly cuts, and the splice between the supervisor's two tables measured quarter by quarter.
- A financial-inclusion index by dimension and annual panels at department and municipality level, with an explicit battery against spurious correlation. Public logbook of mistakes and wins.

Stack: dbt · DuckDB · Python · Parquet · Panel data · Quarto

Skills

Finance domain: FP&A · Close & forecast · SG&A · Revenue recognition (SOX) · Treasury · FX analysis · Budgeting · Constant currency · SAP · JD Edwards · MicroStrategy

Data & engineering: Analytical SQL · PostgreSQL · DuckDB · dbt · Medallion architecture · Python (pandas) · Incremental ETL · Prefect · GitHub Actions · Supabase · Power BI · DAX · Power Query · Tableau · Git · Next.js · Vercel · Panel econometrics · Fixed effects · Causal inference · Time series · Credit scorecards (WoE) · LightGBM · MLflow · ONNX · MLOps · Data testing

Technical stack

- **Financial analysis and FP&A** — SG&A close and forecast across 15+ countries: variance against plan and in constant currency
- **SAP** — Revenue recognition under SOX and automated bank reconciliation, at SLB
- **Excel and financial modelling** — The close and forecast for those 15+ countries, at Neoris EPAM
- **Power BI** — [Seven-table semantic model in TMDL with 17 DAX measures](#), versioned as text
- **Tableau** — [Dashboard that won the BodyTech Trends Hackathon](#), public
- **Python** — Incremental ingestion, backtesting engine, FX decomposition
- **SQL · PostgreSQL** — Three-layer medallion warehouse, more than 58,000 candles in production
- **dbt** — 89 quality tests that run before a single figure is published
- **DuckDB** — Dimensional warehouse over 19 public sources resolved to municipal codes
- **Git · GitHub Actions** — Daily cron in operation, with a rate-limit circuit breaker
- **Panel econometrics** — Two-way fixed effects over 228 observations, wild cluster bootstrap and permutation placebo
- **Diagnostics and robustness** — Driscoll-Kraay for measured cross-sectional dependence, CCE, SLX and four designs against exogeneity
- **Machine learning** — LightGBM over 1.96M SBA loans and 62.4M HMDA applications: AUC 0.7005, +0.0311 over the interpretable scorecard
- **Credit risk and scorecards** — WoE binning with optbinning, out-of-time validation across the COVID shock, calibration error 0.0107
- **Model governance (MLOps)** — Ten promotion gates judged by exit code; one blocks my own model

Education

M.Sc. in Economics

2025 — 2026

Pontificia Universidad Javeriana · thesis filed

Thesis filed in August 2026; graduation expected in November 2026. It is the origin of the financial-inclusion project I keep building in public.

B.Sc. in Economics

2020 — 2024

Pontificia Universidad Javeriana · completed

Systems Technician

2018

SENA · completed

Certifications: Machine Learning Specialization (Stanford · Coursera, 2024) · Cybersecurity Certificate (Google · Coursera, 2024) · Data Science with Python (Platzi, 2023)

Recognition: [Winner — BodyTech Trends Hackathon \(2024\)](#) · Ecopetrol Scholar — Mario Galán Gómez Program (2018)

Remote-ready

4 employers, all remote or hybrid, with teams distributed across 15+ countries. · GMT-5 (Bogotá): full overlap with US and Canadian hours. · Experience working with teams in North America, Argentina and Brazil. · Native Spanish · English B2 · Portuguese A2.