

Davirson Novoa Ramírez

Finance Data Analyst · Financial BI Analyst · Analytics Engineer

I find the variance, I trace it to root cause, I leave the control running itself, and I explain it where the decision gets made. All four, with no translator between the business and the data.

Bogotá, Colombia · GMT-5 · Fully remote · Daily working English with US and European teams · Native Spanish
davirson@davirson.com · linkedin.com/in/davirson-novoa · github.com/DavinsonR · davirson.com

Profile

Finance Data Analyst who owns the SG&A close and forecast process for 15+ countries across the Americas and Europe, from HQ. An economist trained in treasury, SOX billing and FP&A, who builds the data the business runs on instead of waiting for it: Python, SQL and Power BI on real operations. Alongside that I run a production PostgreSQL warehouse in medallion architecture, with dbt transformations, 89 automated quality tests and a Power BI semantic model that refreshes itself daily, and I deployed a credit decisioning system over 95 million records with ten model-governance gates. All of it verifiable: every figure on my site links to the repository that produces it. I bring financial judgment and data engineering in a single hire: no translator, no handoff meeting.

The crossover role

I charge for what I have done in every finance role: in all of them I ended up building the same thing — code and data — because the report the business needed did not exist and nobody else was going to build it. An FX mismatch analysis in Python nobody had asked for, as an intern. Automation that gave every treasury analyst about 10 hours a month back — close to 60 a month across the team. Power BI models that turned a monthly close into a decision. The edge is not knowing Python; plenty of people know Python. It is knowing which question is worth answering before writing the first line. A Finance Data Analyst is neither a data analyst who picked up finance nor a finance person who picked up code — it is the one who needs no translator between them. Hiring that person is two hires in one, without the handoff meeting between them.

Experience

Neoris EPAM

Business Consultant, FP&A

Bogotá, Colombia · remote

Mar 2026 — present

- Own the SG&A close and forecast process for 15+ countries across the Americas and the Eastern Hemisphere (India, Hungary, Spain, Czech Republic), from HQ.
- Accountable for the monthly close and company-wide forecast confirmation; I explain every variance against plan and prior forecast.
- Isolate the FX effect in constant currency and per-currency inflation, so the operating variance reads clean.
- Built the FP&A reporting on advanced Excel, Power Query, Power BI, MicroStrategy, JD Edwards and SAP.

SLB

Bogotá, Colombia · hybrid

Two promotions in 26 months: from treasury to billing specialist for Argentina and Brazil, with an intern reporting to me.

Billing Specialist (Argentina & Brazil)

Dec 2024 — Mar 2026

- Owned revenue recognition in SAP for Argentina and Brazil under Sarbanes-Oxley (SOX) controls.
- Supervised and trained an intern on the billing team.
- Built Power BI models and visualizations for revenue and trend analysis.
- Audited integrated technology projects, with end-to-end traceability and compliance.

Treasury Analyst

Oct 2024 — Dec 2024

- Gave the team back about 60 hours a month — roughly 10 per analyst — by automating manual treasury reporting.
- Automated bank reconciliation in SAP and moved it into Power Automate workflows.

Treasury Intern

Jan 2024 — Jun 2024

- Detected and quantified FX mismatches across Latin American markets in Python, identifying financial risks and hedging opportunities.
- Created Power BI dashboards for cash-flow projections across 15+ countries.
- Automated bank transaction reporting across 10+ regional operations.

LEE Javeriana

Research Assistant

remote

Jul 2023 — Dec 2023

- Modelled Colombian higher-education data for impact-evaluation studies.
- Published the research findings in Power BI dashboards.

Projects in operation

Open source and checkable. None of this is a course exercise.

credit-risk-mlops

Credit decisioning system with model governance — design and build

2026 — deployed, scoring live
github.com/DavinsonR/credit-risk-mlops

- Declining the riskiest 10% of the test portfolio would have avoided \$276.3M in charge-offs — 2.15x what random declines achieve — and the report ships the counterweight too: \$1.99B in good lending forgone.
- Deployed model at AUC 0.7005 over 1.96M SBA 7(a) loans and 62.4M HMDA applications: +0.0311 over an interpretable WoE scorecard, with calibration error 0.0107. The event study runs separately, over 93.4M applications.
- The first reading was 0.9461, and it was not a model but a leak: TermInMonths is overwritten when a loan is liquidated, so the field carried the outcome. Removing it drops the ablation to 0.6621 — and that is the number published.
- Ten promotion gates judged by exit code, one of them designed to block my own model when it misses the threshold.
- Monitoring caught the source changing its vocabulary without notice: a field stopped meaning the same thing between vintages, and nobody would have known without the control.
- The production ONNX artifact scores in the reader's own browser — no server, and no data leaves the page.

Stack: Python · LightGBM · scikit-learn · optbinning (WoE) · MLflow · ONNX · DuckDB · pandera · Docker · Power BI

market-data-medallion

Market data platform — designed, built and operated

2026 — in daily operation
github.com/DavinsonR/market-data-medallion

- Watermark-based incremental ingestion from 4 sources (Tiingo, Coinbase, Kraken, FX) into a three-layer PostgreSQL warehouse: 48 assets and more than 58,000 candles, idempotent and resumable.
- 89 automated dbt data-quality tests plus 171 Python unit tests; the pipeline fails before publishing bad data, not after.
- No-look-ahead backtesting engine with 70/30 out-of-sample validation over 1,392 strategy variants: fewer than one in six in-sample winners survived the blind window.
- FX decomposition for Latin American ADRs, separating company performance from the currency move through the identity $(1+r_USD) \times (1+r_FX) = (1+r_local)$.
- Daily orchestration on GitHub Actions with a rate-limit circuit breaker, running on zero-cost infrastructure.

Stack: PostgreSQL · Medallion architecture · dbt · Python · Prefect · GitHub Actions · Supabase · Power BI (TMDL / PBIP) · Next.js

JARVIS — multi-user product with sensitive data

Own product — designed, built and operated

2026 — in production · open demo
[public demo, no account](#)

- Multi-user model on Postgres with a row policy on 34 tables: the conversion required rebuilding 22 primary keys that identified rows without the user, and pushing the user inside the foreign keys so isolation did not rest on the policy alone.
- Eight CI gates judged by exit code: types, format, lint, 526 tests, build, dependency audit, a guard that hunts personal data inside the built artifact, and a smoke test requiring the anonymous role to be denied across all 55 tables and views.
- A public demo running the five real screens over 120 seed-generated days that never queries the database; a test walks the runtime import graph and fails if any reachable module gets to the data client.
- Day logging in one form and one button, with eight independent writes, an on-device draft and send queue to survive losing signal, and a purpose-built number reader because a browser number field was storing expenses divided by a thousand.

Stack: Next.js 16 · React 19 · TypeScript · Postgres (Supabase) · RLS · Tailwind · PWA · Vitest

Research

My own research, built in public on open data.

Financial inclusion and regional growth in Colombia

M.Sc. in Economics thesis, Javeriana — reproducible research end to end

2026 — thesis filed
[repository on GitHub](#)

- Purpose-built open-data downloader with a verifiable manifest: nineteen sources (financial supervisor, statistics office, ICT and education ministries, national map), row counts verified against the source, metadata-driven typing and a sha256 per file.
- Dimensional warehouse in dbt on DuckDB: department, municipality and period dimensions, financial inclusion in long form by product block, and tests that reconcile each department total with the sum of its municipalities.
- Geographic keys resolved to the municipal code with 100% coverage across the 34 quarterly cuts, and the splice between the supervisor's two tables measured quarter by quarter.
- A financial-inclusion index by dimension and annual panels at department and municipality level, with an explicit battery against spurious correlation. Public logbook of mistakes and wins.

Stack: dbt · DuckDB · Python · Parquet · Panel data · Quarto

Skills

Finance domain: FP&A · Close & forecast · SG&A · Revenue recognition (SOX) · Treasury · FX analysis · Budgeting · Constant currency · SAP · JD Edwards · MicroStrategy · Variance analysis · Month-end close · Cash-flow forecasting · Management reporting · Account reconciliation

Data & engineering: Analytical SQL · PostgreSQL · DuckDB · dbt · Medallion architecture · Python (pandas) · Incremental ETL · Prefect · GitHub Actions · Supabase · Power BI · DAX · Power Query · Tableau · Git · Next.js · Vercel · Panel econometrics · Fixed effects · Causal inference · Time series · Credit scorecards (WoE) · LightGBM · MLflow · ONNX · MLOps · Data testing · Data warehousing · Dimensional modelling (star schema) · ELT · CI/CD · Dashboarding

Technical stack

- **Financial analysis and FP&A** — SG&A close and forecast across 15+ countries: variance against plan and in constant currency
- **SAP** — Revenue recognition under SOX and automated bank reconciliation, at SLB
- **Excel and financial modelling** — The close and forecast for those 15+ countries, at Neoris EPAM
- **Power BI** — [Seven-table semantic model in TMDL with 17 DAX measures, versioned as text](#)
- **Tableau** — [Dashboard that won the BodyTech Trends Hackathon, public](#)
- **Python** — Incremental ingestion, backtesting engine, FX decomposition
- **SQL · PostgreSQL** — Three-layer medallion warehouse, more than 58,000 candles in production
- **dbt** — 89 quality tests that run before a single figure is published
- **DuckDB** — Dimensional warehouse over 19 public sources resolved to municipal codes
- **Git · GitHub Actions** — Daily cron in operation, with a rate-limit circuit breaker
- **Panel econometrics** — Two-way fixed effects over 228 observations, wild cluster bootstrap and permutation placebo
- **Diagnostics and robustness** — Driscoll-Kraay for measured cross-sectional dependence, CCE, SLX and four designs against exogeneity
- **Machine learning** — LightGBM over 1.96M SBA loans and 62.4M HMDA applications: AUC 0.7005, +0.0311 over the interpretable scorecard
- **Credit risk and scorecards** — WoE binning with optbinning, out-of-time validation across the COVID shock, calibration error 0.0107
- **Model governance (MLOps)** — Ten promotion gates judged by exit code; one blocks my own model

Education

M.Sc. in Economics

2025 — 2026

Pontificia Universidad Javeriana · thesis filed

Thesis filed in August 2026; graduation expected in November 2026. It is the origin of the financial-inclusion project I keep building in public.

B.Sc. in Economics

2020 — 2024

Pontificia Universidad Javeriana · completed

Certifications: Machine Learning Specialization (Stanford · Coursera, 2024) · Cybersecurity Certificate (Google · Coursera, 2024)

Recognition: [Winner — BodyTech Trends Hackathon \(2024\)](#) · Ecopetrol Scholar — Mario Galán Gómez Program (2018)

Remote-ready

My entire track record is remote or hybrid, with teams distributed across 15+ countries. · GMT-5 (Bogotá): your US and Canadian hours, in full. · Daily working English with teams in the US and Europe (Hungary, Czech Republic, Spain), plus Argentina and Brazil. · Native Spanish · Portuguese A2.